

# Air Pollution Analysis in Seoul, South Korea

Frantz Alexander

## Case Study:

Analyze and predict the air quality in the Gangnam downtown shopping district for the first of November 2019.

## Libraries

```
In [1]: import time

import warnings

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

from sklearn.metrics import mean_absolute_error
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA

%matplotlib inline
```

```
warnings.filterwarnings("ignore")
```

## Prepare Dataset

### Import Dataset

```
In [2]: def wrangle(filepath, resample_rule = "1H"):
        # Read data into DataFrame
        dataset = pd.read_csv(filepath)

        # Subset station in Gangnam
        station_mask = dataset["Station code"] == 123

        # Subset readings to pm2.5
        pm_mask = dataset["Item code"] == 9

        # subset
        dataset = dataset[station_mask & pm_mask]

        # Drop station code, item code, instrument status
        dataset = (
            dataset.drop(
                columns = ["Station code", "Item code", "Instrument status"]
            )
        )

        # Rename Values to PM2.5 readings
        dataset = dataset.rename(columns = {"Average value": "PM2.5 Readings"})

        # Setting the index to the measurement date
        dataset = dataset.set_index("Measurement date")

        # Setting index to datetime
        dataset.index = pd.to_datetime(dataset.index)

        # Converting index to local time in Seoul
```

```

dataset.index = dataset.index.tz_localize("UTC").tz_convert("Asia/Seoul")

# Remove outliers
outlier_mask_high = dataset["PM2.5 Readings"] < 500

outlier_mask_low = dataset["PM2.5 Readings"] > 0

dataset = dataset[outlier_mask_high & outlier_mask_low]

# Resample and forward-fill
y = dataset["PM2.5 Readings"].resample(resample_rule).mean().fillna(method = "ffill")

return y

```

```
In [3]: y = wrangle("Measurement_info.csv")
```

```
In [4]: y.info()
y.tail(30)
```

```

<class 'pandas.core.series.Series'>
DatetimeIndex: 26280 entries, 2017-01-01 09:00:00+09:00 to 2020-01-01 08:00:00+09:00
Freq: H
Series name: PM2.5 Readings
Non-Null Count  Dtype
-----
26280 non-null  float64
dtypes: float64(1)
memory usage: 410.6 KB

```

```
Out[4]: Measurement date
2019-12-31 03:00:00+09:00    31.0
2019-12-31 04:00:00+09:00    23.0
2019-12-31 05:00:00+09:00    22.0
2019-12-31 06:00:00+09:00    16.0
2019-12-31 07:00:00+09:00    14.0
2019-12-31 08:00:00+09:00    14.0
2019-12-31 09:00:00+09:00    20.0
2019-12-31 10:00:00+09:00    18.0
2019-12-31 11:00:00+09:00    19.0
2019-12-31 12:00:00+09:00    23.0
2019-12-31 13:00:00+09:00    23.0
2019-12-31 14:00:00+09:00    18.0
2019-12-31 15:00:00+09:00    14.0
2019-12-31 16:00:00+09:00    12.0
2019-12-31 17:00:00+09:00    12.0
2019-12-31 18:00:00+09:00    10.0
2019-12-31 19:00:00+09:00     9.0
2019-12-31 20:00:00+09:00     7.0
2019-12-31 21:00:00+09:00    11.0
2019-12-31 22:00:00+09:00    12.0
2019-12-31 23:00:00+09:00     9.0
2020-01-01 00:00:00+09:00    13.0
2020-01-01 01:00:00+09:00    16.0
2020-01-01 02:00:00+09:00    14.0
2020-01-01 03:00:00+09:00    15.0
2020-01-01 04:00:00+09:00    15.0
2020-01-01 05:00:00+09:00    16.0
2020-01-01 06:00:00+09:00    14.0
2020-01-01 07:00:00+09:00    15.0
2020-01-01 08:00:00+09:00    13.0
Freq: H, Name: PM2.5 Readings, dtype: float64
```

## Explore Data

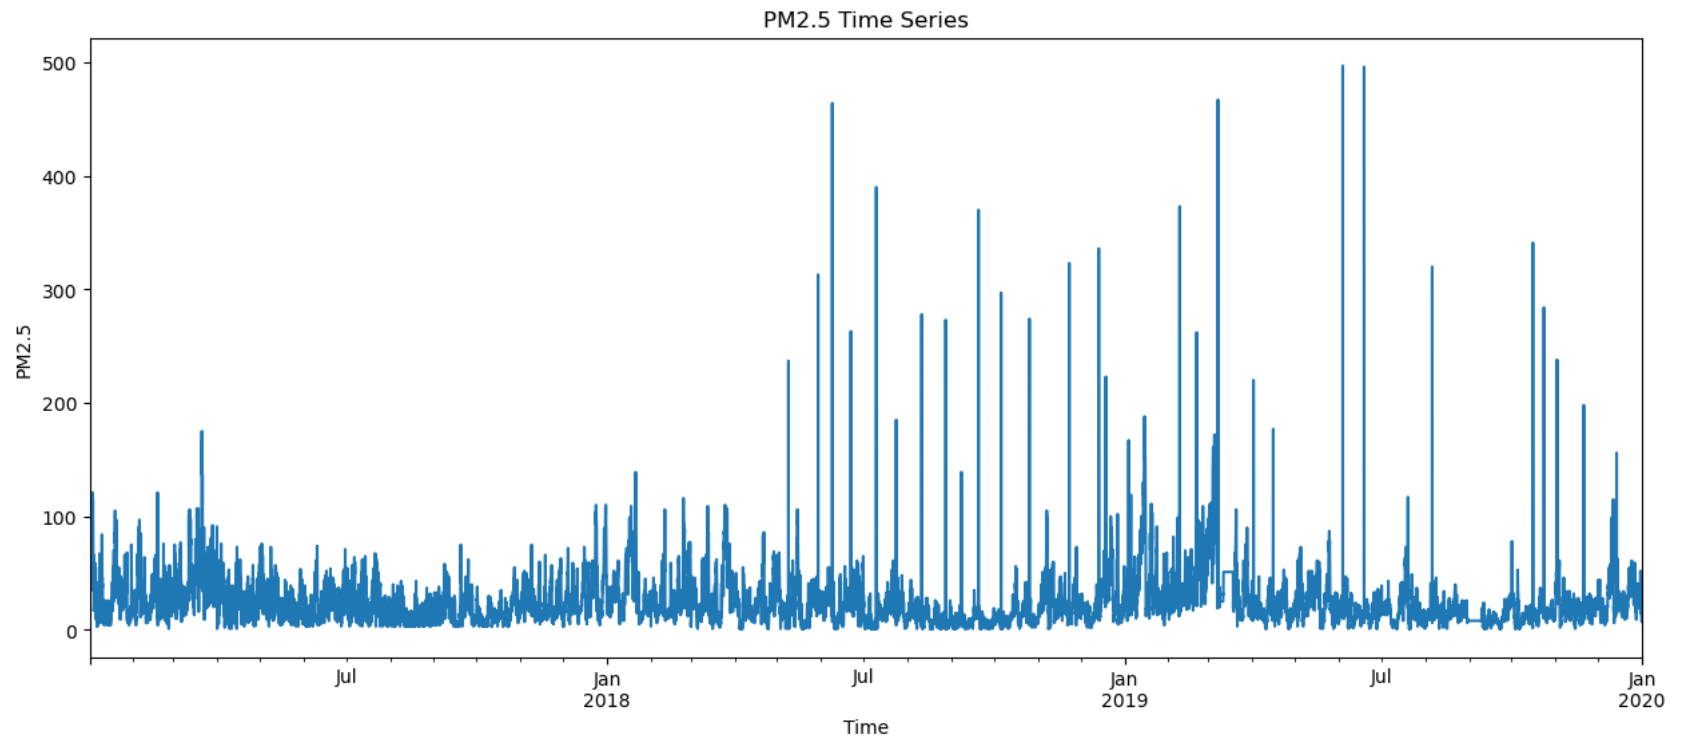


```
In [5]: y.describe()
```

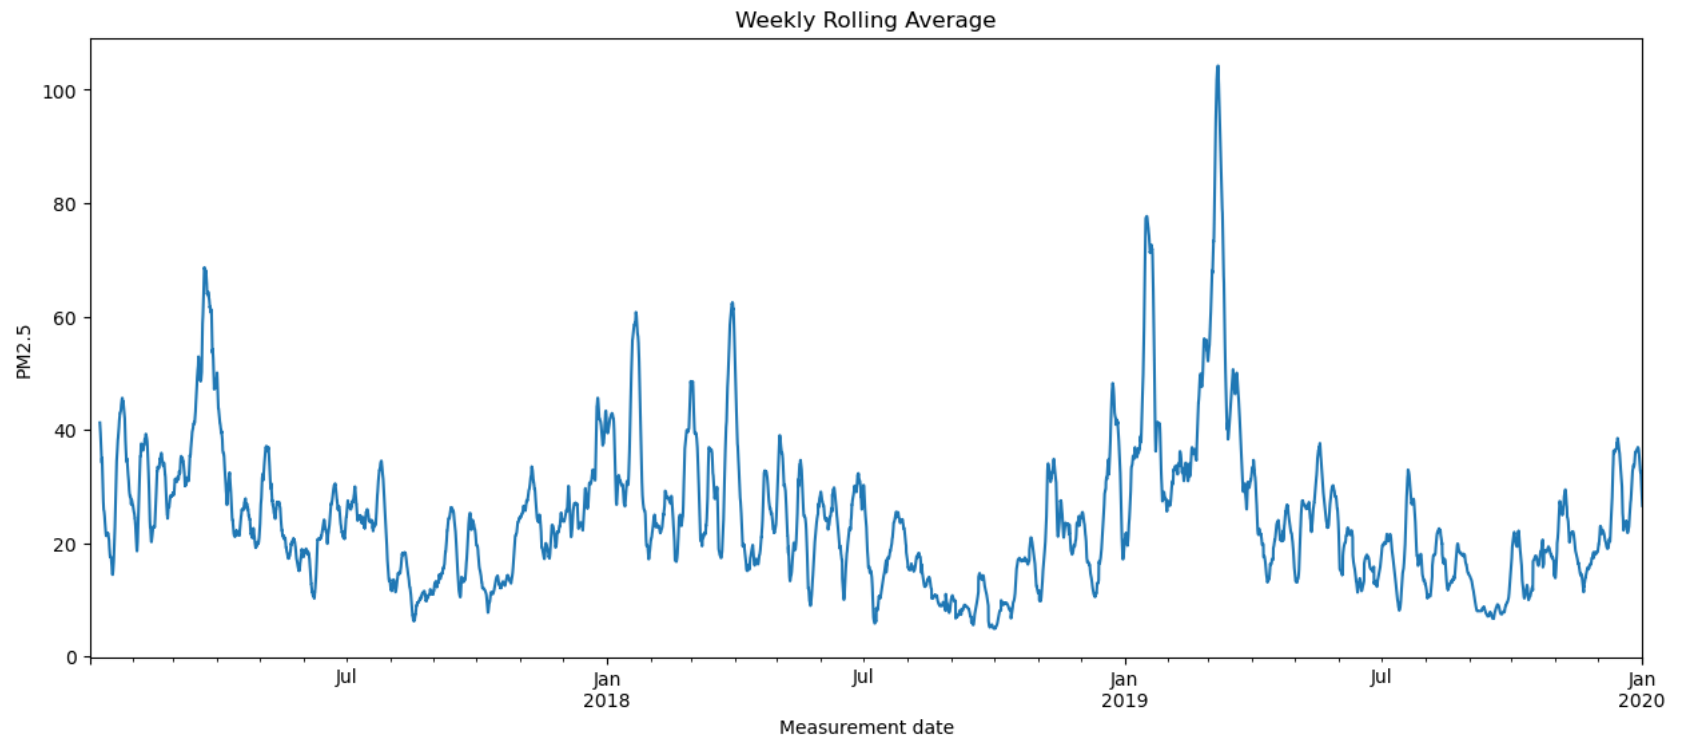
```
Out[5]: count    26280.000000
        mean      24.465715
        std       22.249360
        min        1.000000
        25%       11.000000
        50%       19.000000
        75%       31.000000
        max       497.000000
        Name: PM2.5 Readings, dtype: float64
```

We removed readings that were above 500 as they were highly likely to be measurement errors. Readings that were below 0 were also removed.

```
In [6]: fig, ax = plt.subplots(figsize = (15, 6))
        y.plot(
            xlabel = "Time",
            ylabel = "PM2.5",
            title = "PM2.5 Time Series",
            ax = ax
        );
```



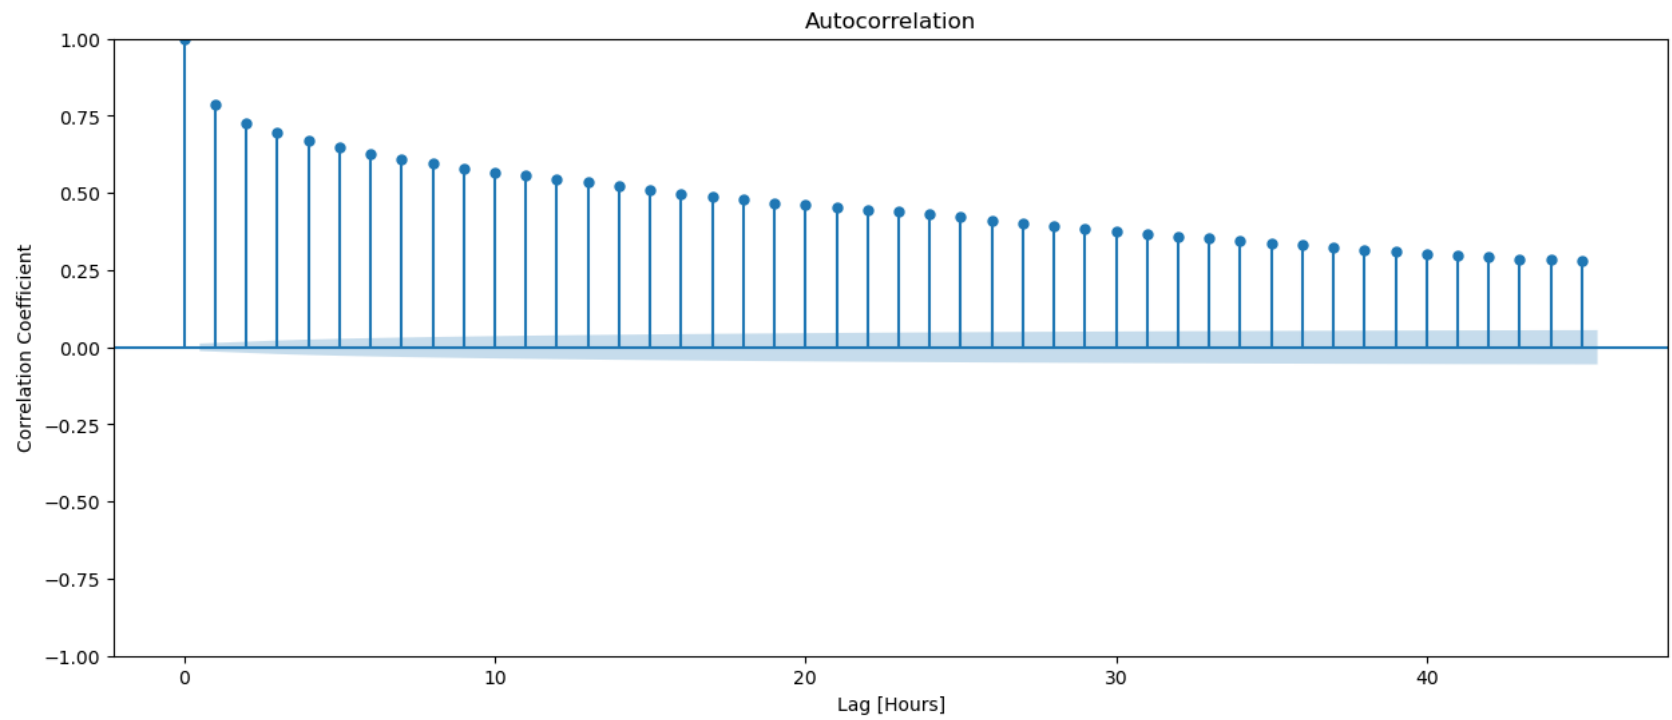
```
In [7]: fig, ax = plt.subplots(figsize = (15, 6))
y.rolling(168).mean().plot(
    ax = ax,
    ylabel = "PM2.5",
    title = "Weekly Rolling Average"
);
```



Using a rolling average smooths out the data and aids in identifying general trends that would be important to model.

Create an ACF Plot for the data in `y`.

```
In [8]: fig, ax = plt.subplots(figsize = (15, 6))
plot_acf(y, ax = ax)
plt.xlabel("Lag [Hours]")
plt.ylabel("Correlation Coefficient");
```



The ACF plot illustrates how the Correlation Coefficient changes as the time lag increases.

The blue band at the bottom of the graph shows that the measurements within it are statistically insignificant.

That is because they are very close to a Correlation Coefficient of 0.

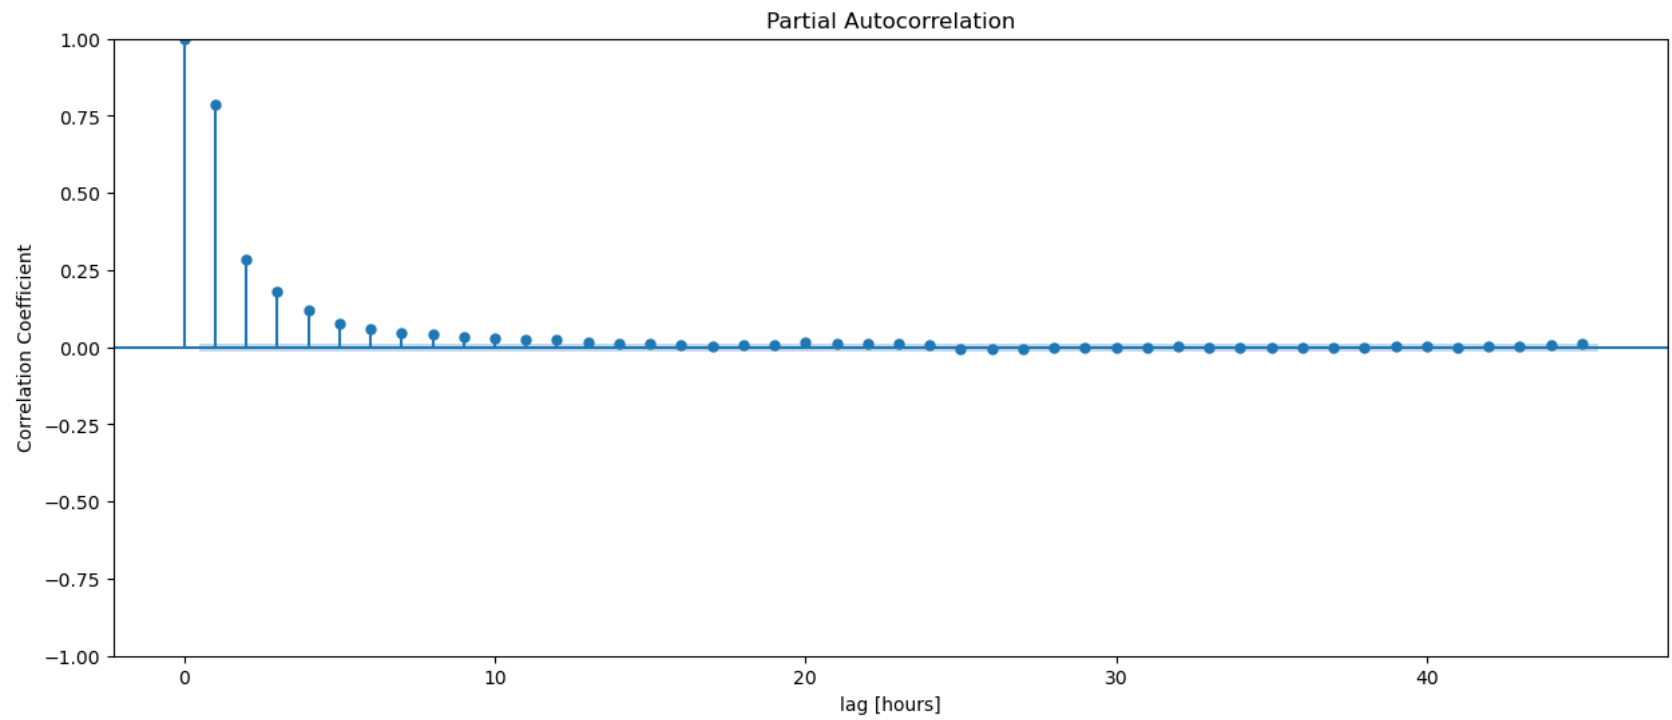
We are only interested in measurements that are far above or below the band.

We use the Partial AutoCorrelation Function plot(PACF) to determine how many lag terms we want in our model.

Create PACF Plot

```
In [9]: fig, ax = plt.subplots(figsize = (15, 6))
plot_pacf(y, ax = ax)
plt.xlabel("lag [hours]")
plt.ylabel("Correlation Coefficient");
```





Once you pull out the lag at 1 hour, the predictive power drastically drops as time passes. The predictive power continues to drop until a time lag of 12. Then again at lags 22, and 23 hours.

At lag 26, the correlation coefficient becomes slightly negative until a lag of 45.

Adding more time lags beyond 45 would not benefit our model.

## Split Data



Training Set

```
In [10]: y["2019-10-01":"2019-10-31"]
```

```
Out[10]: Measurement date
2019-10-01 00:00:00+09:00    25.0
2019-10-01 01:00:00+09:00    22.0
2019-10-01 02:00:00+09:00    39.0
2019-10-01 03:00:00+09:00    57.0
2019-10-01 04:00:00+09:00    57.0
...
2019-10-31 19:00:00+09:00    12.0
2019-10-31 20:00:00+09:00    12.0
2019-10-31 21:00:00+09:00    16.0
2019-10-31 22:00:00+09:00    24.0
2019-10-31 23:00:00+09:00    25.0
Freq: H, Name: PM2.5 Readings, Length: 744, dtype: float64

Test set
```

```
In [11]: y["2019-11-01":"2019-11-30"]
```

```
Out[11]: Measurement date
2019-11-01 00:00:00+09:00    30.0
2019-11-01 01:00:00+09:00    37.0
2019-11-01 02:00:00+09:00    37.0
2019-11-01 03:00:00+09:00    40.0
2019-11-01 04:00:00+09:00    40.0
...
2019-11-30 19:00:00+09:00    28.0
2019-11-30 20:00:00+09:00    24.0
2019-11-30 21:00:00+09:00    25.0
2019-11-30 22:00:00+09:00    26.0
2019-11-30 23:00:00+09:00    26.0
Freq: H, Name: PM2.5 Readings, Length: 720, dtype: float64
```

Create Training Variable and Test Variable

```
In [12]: y_train = y["2019-10-01":"2019-10-31"]
y_test = y["2019-11-01"]
```

```
In [13]: print(len(y_train))
print(len(y_test))
```

# Model Building

## Baseline

Calculate the baseline Mean Absolute Error

```
In [14]: y_train_mean = y_train.mean()
y_pred_baseline = [y_train_mean]*len(y_train)
mae_baseline = mean_absolute_error(y_train, y_pred_baseline)

print("Mean P2 Reading:", round(y_train_mean, 2))
print("Baseline MAE", round(mae_baseline, 2))
```

```
Mean P2 Reading: 15.98
Baseline MAE 8.08
```

## Iterate

### Hyper Parameters

```
In [15]: p_params = range(0, 46, 5)
q_params = range(0, 4, 1)
```

```
In [16]: list(p_params)
```

```
Out[16]: [0, 5, 10, 15, 20, 25, 30, 35, 40, 45]
```

```
In [17]: list(q_params)
```

```
Out[17]: [0, 1, 2, 3]
```

```
In [18]: # Create an empty dictionary for MAE values
mae_grid = {}
for p in p_params:
    # Create new key in dict with empty List
    mae_grid[p] = []
    for q in q_params:
        order = (p, 0, q)
        # Start Timing
        start_time = time.time()
        model = ARIMA(y_train, order = order).fit()
        # Calculate the Elapsed Time
        elapsed_time = round(time.time() - start_time, 2)
        print(f"Trained ARIMA model {order} in {elapsed_time} seconds.")
        # Generate in-sample predictions
        y_pred = model.predict()
        # Calculate training MAE
        mae = mean_absolute_error(y_train, y_pred)
        print(mae)
        # Add MAE to Dictionary
        mae_grid[p].append(mae)
```

Trained ARIMA model (0, 0, 0) in 0.16 seconds.  
8.081945672169077  
Trained ARIMA model (0, 0, 1) in 0.12 seconds.  
5.469754371426196  
Trained ARIMA model (0, 0, 2) in 0.17 seconds.  
4.545075159100064  
Trained ARIMA model (0, 0, 3) in 0.27 seconds.  
4.4256982832994  
Trained ARIMA model (5, 0, 0) in 0.18 seconds.  
4.226508439164455  
Trained ARIMA model (5, 0, 1) in 0.71 seconds.  
4.231962358179267  
Trained ARIMA model (5, 0, 2) in 1.33 seconds.  
4.175676546497824  
Trained ARIMA model (5, 0, 3) in 1.4 seconds.  
4.145377676127871  
Trained ARIMA model (10, 0, 0) in 0.67 seconds.  
4.165848676340261  
Trained ARIMA model (10, 0, 1) in 2.52 seconds.  
4.186803316668656  
Trained ARIMA model (10, 0, 2) in 2.2 seconds.  
4.142727910159227  
Trained ARIMA model (10, 0, 3) in 1.42 seconds.  
4.137354466464098  
Trained ARIMA model (15, 0, 0) in 1.15 seconds.  
4.1406348006994005  
Trained ARIMA model (15, 0, 1) in 1.02 seconds.  
4.140808825169648  
Trained ARIMA model (15, 0, 2) in 1.23 seconds.  
4.141042480804938  
Trained ARIMA model (15, 0, 3) in 2.0 seconds.  
4.155261991090604  
Trained ARIMA model (20, 0, 0) in 1.76 seconds.  
4.137845619147744  
Trained ARIMA model (20, 0, 1) in 1.58 seconds.  
4.139377925758959  
Trained ARIMA model (20, 0, 2) in 1.85 seconds.  
4.139552909198084  
Trained ARIMA model (20, 0, 3) in 1.65 seconds.  
4.139419369988037  
Trained ARIMA model (25, 0, 0) in 3.1 seconds.  
4.1289408575348

Trained ARIMA model (25, 0, 1) in 2.38 seconds.  
4.12886561602995  
Trained ARIMA model (25, 0, 2) in 5.38 seconds.  
4.162416165626995  
Trained ARIMA model (25, 0, 3) in 2.85 seconds.  
4.1295606686006225  
Trained ARIMA model (30, 0, 0) in 5.23 seconds.  
4.127449487856752  
Trained ARIMA model (30, 0, 1) in 3.49 seconds.  
4.127925492781879  
Trained ARIMA model (30, 0, 2) in 4.27 seconds.  
4.127990474157672  
Trained ARIMA model (30, 0, 3) in 10.65 seconds.  
4.1524915306381  
Trained ARIMA model (35, 0, 0) in 7.59 seconds.  
4.134496001416412  
Trained ARIMA model (35, 0, 1) in 6.46 seconds.  
4.135085438395935  
Trained ARIMA model (35, 0, 2) in 7.0 seconds.  
4.135119956234608  
Trained ARIMA model (35, 0, 3) in 21.15 seconds.  
4.1604595502463235  
Trained ARIMA model (40, 0, 0) in 10.46 seconds.  
4.139770108980159  
Trained ARIMA model (40, 0, 1) in 7.36 seconds.  
4.13874002223658  
Trained ARIMA model (40, 0, 2) in 9.66 seconds.  
4.138679637108064  
Trained ARIMA model (40, 0, 3) in 30.43 seconds.  
4.187551368060375  
Trained ARIMA model (45, 0, 0) in 16.05 seconds.  
4.129178974160771  
Trained ARIMA model (45, 0, 1) in 11.87 seconds.  
4.128304612996378  
Trained ARIMA model (45, 0, 2) in 40.63 seconds.  
4.1580243638611  
Trained ARIMA model (45, 0, 3) in 40.4 seconds.  
4.122860154601666

In [19]: `print(mae_grid)`

```
{0: [8.081945672169077, 5.469754371426196, 4.545075159100064, 4.4256982832994], 5: [4.226508439164455, 4.23196235817
9267, 4.175676546497824, 4.145377676127871], 10: [4.165848676340261, 4.186803316668656, 4.142727910159227, 4.1373544
66464098], 15: [4.1406348006994005, 4.140808825169648, 4.141042480804938, 4.155261991090604], 20: [4.13784561914774
4, 4.139377925758959, 4.139552909198084, 4.139419369988037], 25: [4.1289408575348, 4.12886561602995, 4.1624161656269
95, 4.1295606686006225], 30: [4.127449487856752, 4.127925492781879, 4.127990474157672, 4.1524915306381], 35: [4.1344
96001416412, 4.135085438395935, 4.135119956234608, 4.1604595502463235], 40: [4.139770108980159, 4.13874002223658, 4.
138679637108064, 4.187551368060375], 45: [4.129178974160771, 4.128304612996378, 4.1580243638611, 4.122860154601666]}
```

```
In [20]: def gold_min(xs):
    m = xs.to_numpy().min()
    color = {True: "background-color: #c78f2e", False: ""}
    is_min = (xs == m).replace(color)
    return is_min

mae_df = pd.DataFrame(mae_grid)
mae_df.round(4).style.apply(gold_min, axis = None)
```

```
Out[20]:
```

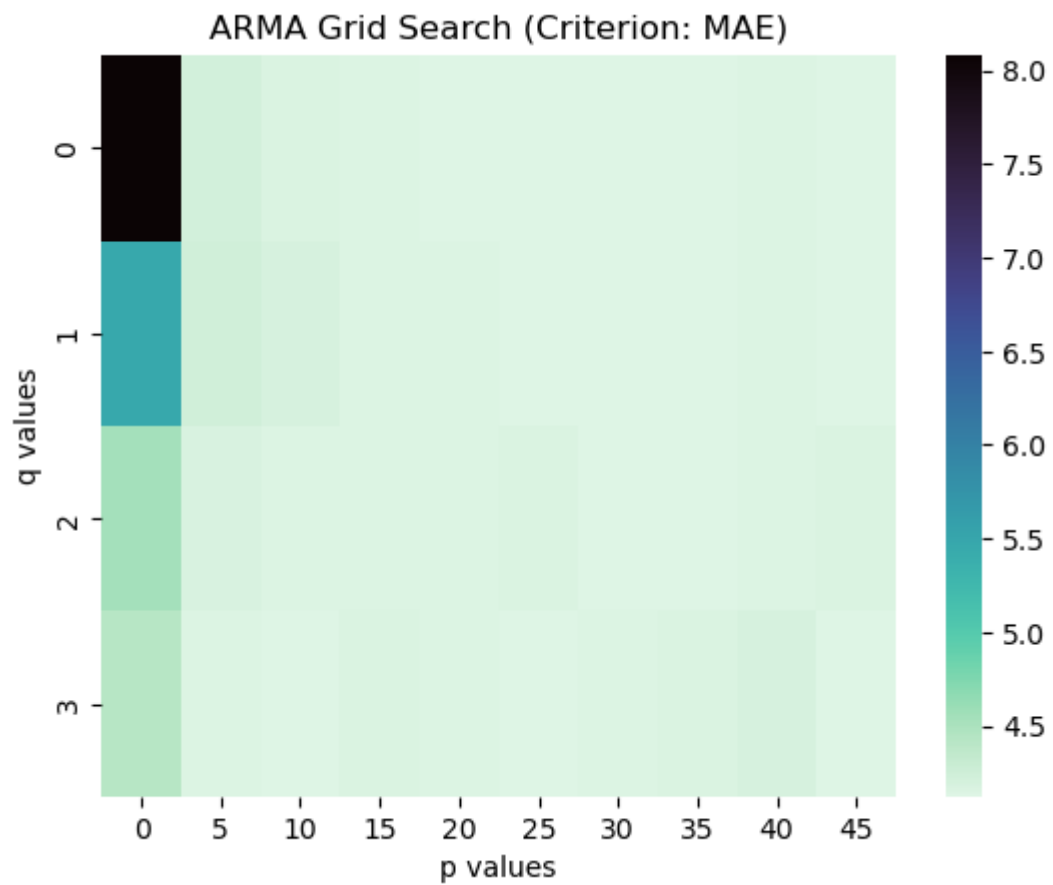
	0	5	10	15	20	25	30	35	40	45
0	8.081900	4.226500	4.165800	4.140600	4.137800	4.128900	4.127400	4.134500	4.139800	4.129200
1	5.469800	4.232000	4.186800	4.140800	4.139400	4.128900	4.127900	4.135100	4.138700	4.128300
2	4.545100	4.175700	4.142700	4.141000	4.139600	4.162400	4.128000	4.135100	4.138700	4.158000
3	4.425700	4.145400	4.137400	4.155300	4.139400	4.129600	4.152500	4.160500	4.187600	4.122900

The optimal the values for p and q are the parameters that best balance model performance with computational time.

In this case, where p = 30 and q = 1.

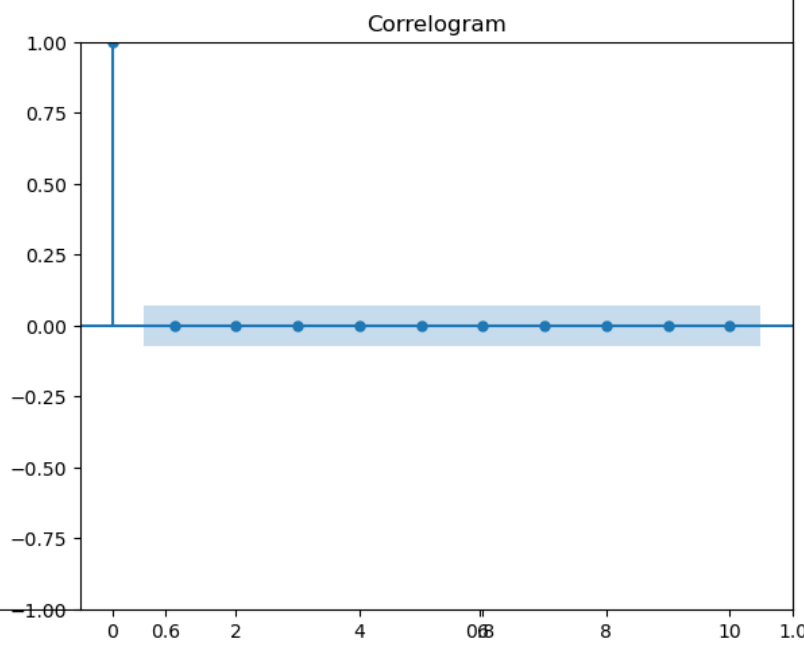
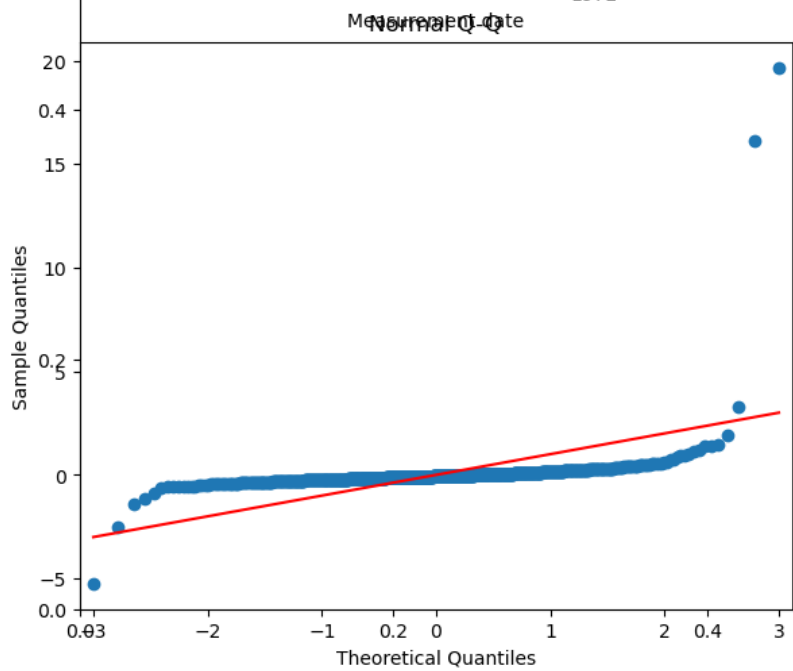
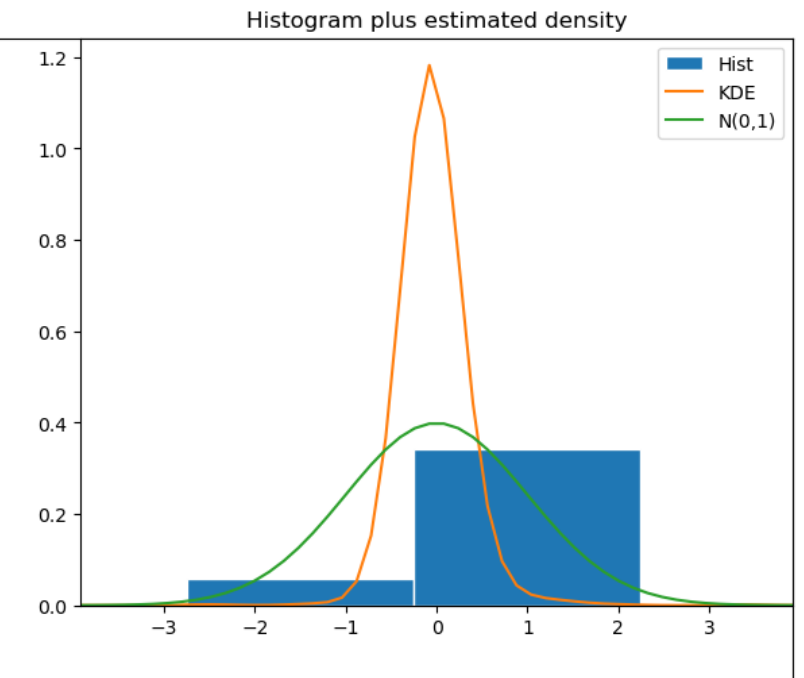
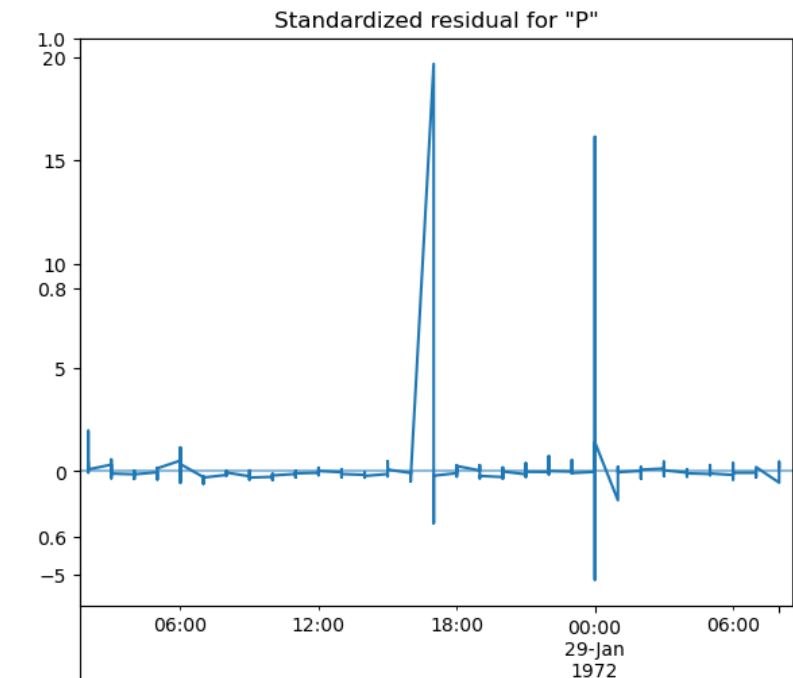
It is important to note that the parameters that produce the best model performance are when p = 30 and q = 0.

```
In [21]: sns.heatmap(
    data = mae_df,
    cmap = "mako_r"
)
plt.xlabel("p values")
plt.ylabel("q values")
plt.title("ARMA Grid Search (Criterion: MAE)");
```



```
In [22]: fig, ax = plt.subplots(figsize = (15, 12))  
         model.plot_diagnostics(fig = fig);
```





## Evaluate



### Walk-Forward Validation

```
In [23]: y_pred_wfv = pd.Series().astype(str)
history = y_train.copy()
for i in range(len(y_test)):
    model = ARIMA(history, order = (30, 0, 1)).fit()
    next_pred = model.forecast()
    y_pred_wfv = pd.concat([y_pred_wfv, next_pred])
    history = pd.concat([history, y_test[next_pred.index]])
```

```
In [24]: test_mae = mean_absolute_error(y_test, y_pred_wfv)
print("Test Mae (Walk-Forward Validation):", round(test_mae, 2))
```

Test Mae (Walk-Forward Validation): 8.76

## Communicate



```
In [25]: df_predictions = pd.DataFrame({"y_test" : y_test, "y_pred_wfv": y_pred_wfv})
fig = px.line(df_predictions, labels = {"value": "PM2.5"})
fig.show()
```

The Baseline Mae was 8.08.

The mae from the Test Data with walk-forward validation was at 8.76.

November 1st experienced significantly higher air pollution than what was represented in the training data from the month of October.

Thus, the model underestimated the pollution increase.

